

Type *Markdown* and LaTeX: α^2

```
In [9]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [10]: df = pd.read_csv(r"E:\154\3_Fitness-1 - 3_Fitness-1.csv")
df
```

Out[10]:

	Row Labels	Sum of Jan	Sum of Feb	Sum of Mar	Sum of Total Sales
0	A	5.62%	7.73%	6.16%	75
1	B	4.21%	17.27%	19.21%	160
2	C	9.83%	11.60%	5.17%	101
3	D	2.81%	21.91%	7.88%	127
4	E	25.28%	10.57%	11.82%	179
5	F	8.15%	16.24%	18.47%	167
6	G	18.54%	8.76%	17.49%	171
7	H	25.56%	5.93%	13.79%	170
8	Grand Total	100.00%	100.00%	100.00%	1150

```
In [11]: df.head()
```

Out[11]:

	Row Labels	Sum of Jan	Sum of Feb	Sum of Mar	Sum of Total Sales
0	A	5.62%	7.73%	6.16%	75
1	B	4.21%	17.27%	19.21%	160
2	C	9.83%	11.60%	5.17%	101
3	D	2.81%	21.91%	7.88%	127
4	E	25.28%	10.57%	11.82%	179

Data cleaning and pre processing

In [12]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9 entries, 0 to 8
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Row Labels            9 non-null     object
1   Sum of Jan             9 non-null     object
2   Sum of Feb             9 non-null     object
3   Sum of Mar             9 non-null     object
4   Sum of Total Sales     9 non-null     int64
dtypes: int64(1), object(4)
memory usage: 488.0+ bytes
```

In [13]: `df.describe()`

Out[13]:

Sum of Total Sales	
count	9.000000
mean	255.555556
std	337.332963
min	75.000000
25%	127.000000
50%	167.000000
75%	171.000000
max	1150.000000


In [14]: `df.columns`

Out[14]: Index(['Row Labels', 'Sum of Jan', 'Sum of Feb', 'Sum of Mar',
'Sum of Total Sales'],
dtype='object')

EDA and VISUALIZATION

In [15]: `sns.pairplot(df)`


Out[15]: <seaborn.axisgrid.PairGrid at 0x201d40f67f0>



```
In [16]: sns.distplot(df["Sum of Total Sales"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[16]: <AxesSubplot:xlabel='Sum of Total Sales', ylabel='Density'>
```



```
In [17]: df1 = df[['Row Labels', 'Sum of Jan', 'Sum of Feb', 'Sum of Mar',  
                'Sum of Total Sales']]
```

```
In [18]: sns.heatmap(df1.corr())
```

```
Out[18]: <AxesSubplot:>
```



```
In [19]: x = df1[['Sum of Total Sales', 'Sum of Total Sales']]  
y = df1['Sum of Total Sales']
```

split the data into training and test data

```
In [20]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
```

```
In [21]: lr = LinearRegression()
lr.fit(x_train, y_train)
```

```
Out[21]: LinearRegression()
```

```
In [22]: lr.intercept_
```

```
Out[22]: 0.0
```


```
In [23]: coeff = pd.DataFrame(lr.coef_, x.columns, columns = ['Co-efficient'])
coeff
```

```
Out[23]:
```

	Co-efficient
Sum of Total Sales	0.5
Sum of Total Sales	0.5

```
In [24]: prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

```
Out[24]: <matplotlib.collections.PathCollection at 0x201d6d95ca0>
```



```
In [25]: lr.score(x_test,y_test)
```

```
Out[25]: 1.0
```

Accuracy

```
In [26]: lr.score(x_train,y_train)
```

```
Out[26]: 1.0
```

```
In [27]: from sklearn.linear_model import Ridge,Lasso
```

```
In [28]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[28]: Ridge(alpha=10)
```

```
In [29]: rr=Ridge(alpha=10)
         rr.fit(x_train,y_train)
```

```
Out[29]: Ridge(alpha=10)
```

```
In [30]: rr.score(x_train,y_train)
```

```
Out[30]: 0.9999999999648033
```

ElasticNet

```
In [31]: from sklearn.linear_model import ElasticNet
         en = ElasticNet()
         en.fit(x_train,y_train)
```

```
Out[31]: ElasticNet()
```

```
In [32]: print(en.coef_)
         [9.99989321e-01 7.11916051e-06]
```

```
In [33]: print(en.intercept_)
         0.0011177320716342365
```

```
In [34]: print(en.predict(x_test))
         [171.00050903  75.00085076 170.00051259]
```

```
In [35]: print(en.score(x_test,y_test))
         0.9999999997951458
```

Evaluation Metrics

```
In [36]: from sklearn import metrics
```

```
In [37]: print("Mean Absolute Error",metrics.mean_absolute_error(y_test,prediction))
         Mean Absolute Error 0.0
```

```
In [38]: print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
         Mean Squared Error: 0.0
```

```
In [39]: print("Root Mean Absolute Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
         Root Mean Absolute Error: 0.0
```

```
In [ ]:
```

In []: