

Type *Markdown* and LaTeX:  $\alpha^2$

```
In [1]: #import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: #import dataset
df=pd.read_csv(r"E:\154\18_world-data-2023.csv").dropna()
df
```

Out[2]:

	Country	Density\n(P/Km2)	Abbreviation	Agricultural Land( %)	Land Area(Km2)	Armed Forces size	Birth Rate	Ca C
0	Afghanistan	60	AF	58.10%	652,230	323,000	32.49	
1	Albania	105	AL	43.10%	28,748	9,000	11.78	3
2	Algeria	18	DZ	17.40%	2,381,741	317,000	24.28	2
4	Angola	26	AO	47.50%	1,246,700	117,000	40.73	2
6	Argentina	17	AR	54.30%	2,780,400	105,000	17.02	
...	...	...	...	...	...	...	...	
185	United Kingdom	281	GB	71.70%	243,610	148,000	11.00	
186	United States	36	US	44.40%	9,833,517	1,359,000	11.60	

In [3]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 110 entries, 0 to 193
Data columns (total 35 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Country                              110 non-null    object
 1   Density\n(P/Km2)                    110 non-null    object
 2   Abbreviation                        110 non-null    object
 3   Agricultural Land( %)              110 non-null    object
 4   Land Area(Km2)                     110 non-null    object
 5   Armed Forces size                  110 non-null    object
 6   Birth Rate                         110 non-null    float64
 7   Calling Code                       110 non-null    float64
 8   Capital/Major City                 110 non-null    object
 9   Co2-Emissions                     110 non-null    object
10   CPI                                110 non-null    object
11   CPI Change (%)                     110 non-null    object
12   Currency-Code                     110 non-null    object
13   ...

```

In [4]: *#to display top 5 rows*  
df.head()

Out[4]:

	Country	Density\n(P/Km2)	Abbreviation	Agricultural Land( %)	Land Area(Km2)	Armed Forces size	Birth Rate	Calling Code	
0	Afghanistan	60	AF	58.10%	652,230	323,000	32.49	93.0	
1	Albania	105	AL	43.10%	28,748	9,000	11.78	355.0	
2	Algeria	18	DZ	17.40%	2,381,741	317,000	24.28	213.0	
4	Angola	26	AO	47.50%	1,246,700	117,000	40.73	244.0	
6	Argentina	17	AR	54.30%	2,780,400	105,000	17.02	54.0	

5 rows × 35 columns

## Data cleaning and Pre-Processing

```
In [5]: #To find null values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 110 entries, 0 to 193
Data columns (total 35 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   Country                                                                110 non-null    object
1   Density (P/Km2)                                                        110 non-null    object
2   Abbreviation                                                            110 non-null    object
3   Agricultural Land( %)                                                  110 non-null    object
4   Land Area(Km2)                                                         110 non-null    object
5   Armed Forces size                                                      110 non-null    object
6   Birth Rate                                                             110 non-null    float64
7   Calling Code                                                           110 non-null    float64
8   Capital/Major City                                                     110 non-null    object
9   Co2-Emissions                                                          110 non-null    object
10  CPI                                                                    110 non-null    object
11  CPI Change (%)                                                         110 non-null    object
12  Currency-Code                                                          110 non-null    object
13  Fertility Rate                                                         110 non-null    float64
14  Forested Area (%)                                                      110 non-null    object
15  Gasoline Price                                                         110 non-null    object
16  GDP                                                                    110 non-null    object
17  Gross primary education enrollment (%) 110 non-null    object
18  Gross tertiary education enrollment (%) 110 non-null    object
19  Infant mortality                                                       110 non-null    float64
20  Largest city                                                            110 non-null    object
21  Life expectancy                                                        110 non-null    float64
22  Maternal mortality ratio                                               110 non-null    float64
23  Minimum wage                                                           110 non-null    object
24  Official language                                                      110 non-null    object
25  Out of pocket health expenditure 110 non-null    object
26  Physicians per thousand                                                110 non-null    float64
27  Population                                                             110 non-null    object
28  Population: Labor force participation (%) 110 non-null    object
29  Tax revenue (%)                                                        110 non-null    object
30  Total tax rate                                                         110 non-null    object
31  Unemployment rate                                                      110 non-null    object
32  Urban_population                                                       110 non-null    object
33  Latitude                                                               110 non-null    float64
34  Longitude                                                              110 non-null    float64
dtypes: float64(9), object(26)
memory usage: 30.9+ KB
```

```
In [6]: # To display summary of statistics
df.describe()
```

```
Out[6]:
```

	Birth Rate	Calling Code	Fertility Rate	Infant mortality	Life expectancy	Maternal mortality ratio	Physicians per thousand	
<b>count</b>	110.000000	110.000000	110.000000	110.000000	110.000000	110.000000	110.000000	11
<b>mean</b>	20.196455	344.290909	2.672182	20.271818	72.671818	137.227273	1.919182	2
<b>std</b>	10.039056	341.231562	1.308142	18.453214	7.000788	201.171462	1.598116	2
<b>min</b>	6.400000	1.000000	0.980000	1.700000	54.300000	2.000000	0.010000	-4
<b>25%</b>	11.075000	70.000000	1.682500	6.100000	67.625000	15.250000	0.467500	
<b>50%</b>	17.830000	239.500000	2.200000	13.600000	74.400000	41.000000	1.640000	2
<b>75%</b>	27.962500	420.750000	3.505000	31.500000	77.350000	176.000000	3.007500	4
<b>max</b>	46.080000	1876.000000	6.910000	78.500000	83.300000	1120.000000	7.120000	6

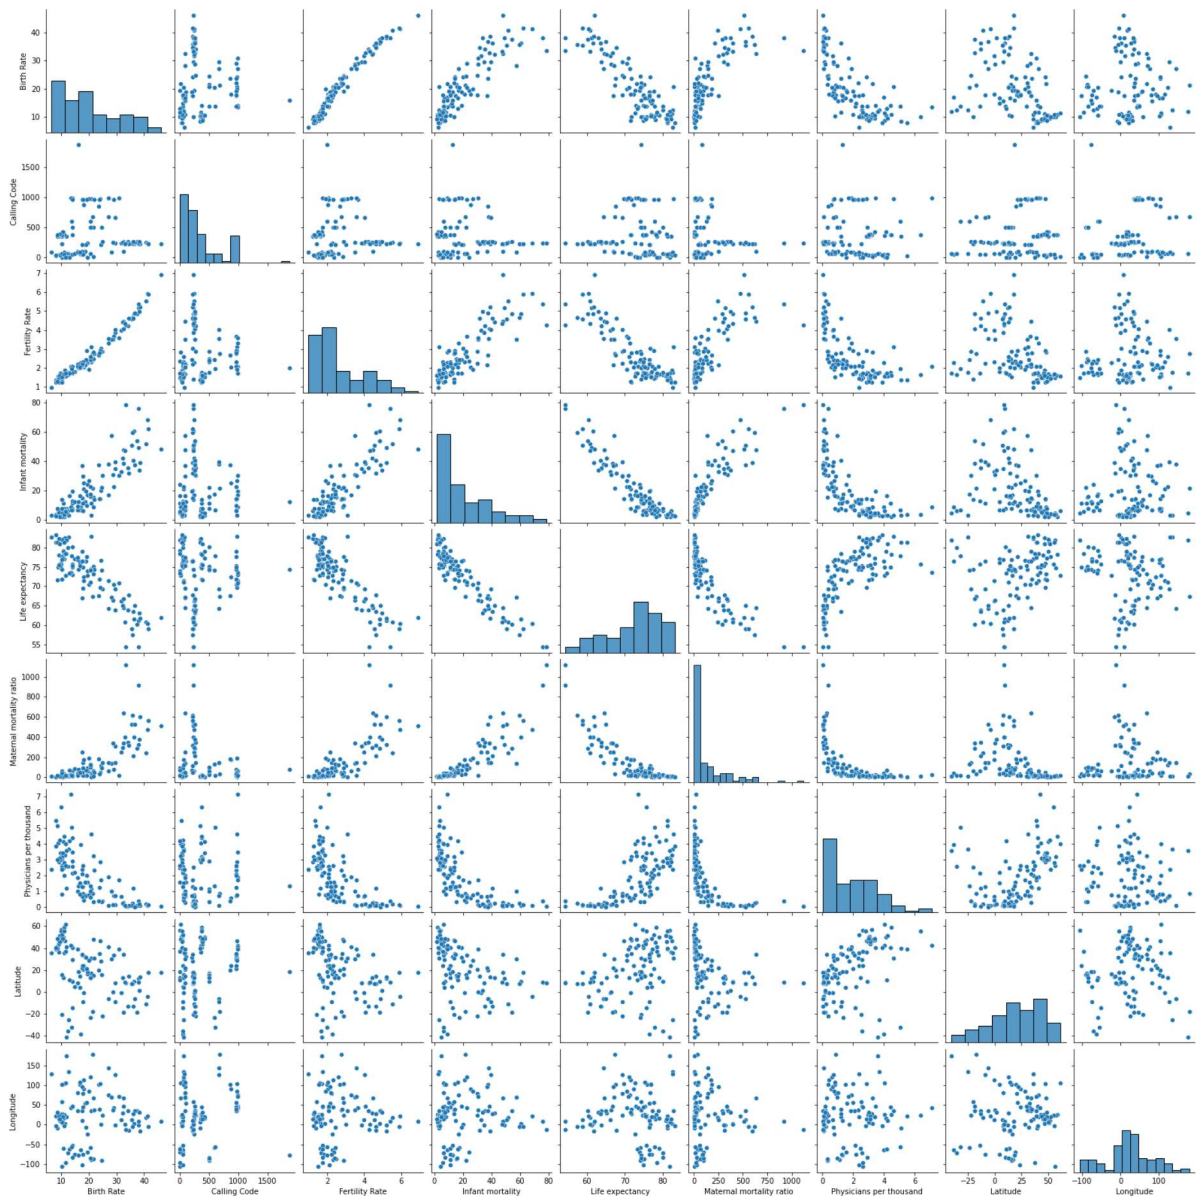
```
In [7]: #To Display column heading
df.columns
```

```
Out[7]: Index(['Country', 'Density\n(P/Km2)', 'Abbreviation', 'Agricultural Land(
%)',
              'Land Area(Km2)', 'Armed Forces size', 'Birth Rate', 'Calling Code',
              'Capital/Major City', 'Co2-Emissions', 'CPI', 'CPI Change (%)',
              'Currency-Code', 'Fertility Rate', 'Forested Area (%)',
              'Gasoline Price', 'GDP', 'Gross primary education enrollment (%)',
              'Gross tertiary education enrollment (%)', 'Infant mortality',
              'Largest city', 'Life expectancy', 'Maternal mortality ratio',
              'Minimum wage', 'Official language', 'Out of pocket health expenditur
e',
              'Physicians per thousand', 'Population',
              'Population: Labor force participation (%)', 'Tax revenue (%)',
              'Total tax rate', 'Unemployment rate', 'Urban_population', 'Latitude',
              'Longitude'],
              dtype='object')
```

## EDA and VISUALIZATION

```
In [8]: sns.pairplot(df)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x1b59f578a00>
```




```
In [9]: sns.distplot(df['Birth Rate'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
Out[9]: <AxesSubplot:xlabel='Birth Rate', ylabel='Density'>
```



```
In [10]: df1=df[['Country', 'Density\n(P/Km2)', 'Abbreviation', 'Agricultural Land( %)',
                'Land Area(Km2)', 'Armed Forces size', 'Birth Rate', 'Calling Code',
                'Capital/Major City', 'Co2-Emissions', 'CPI', 'CPI Change (%)',
                'Currency-Code', 'Fertility Rate', 'Forested Area (%)',
                'Gasoline Price', 'GDP', 'Gross primary education enrollment (%)',
                'Gross tertiary education enrollment (%)', 'Infant mortality',
                'Largest city', 'Life expectancy', 'Maternal mortality ratio',
                'Minimum wage', 'Official language', 'Out of pocket health expenditure',
                'Physicians per thousand', 'Population',
                'Population: Labor force participation (%)', 'Tax revenue (%)',
                'Total tax rate', 'Unemployment rate', 'Urban_population', 'Latitude',
                'Longitude']]
```

## Plot Using Heat Map

```
In [11]: sns.heatmap(df1.corr())
```

```
Out[11]: <AxesSubplot:>
```



## To Train The Model-Model Building

we are going to train Linera Regression Model;We need to split out data into two variables x and y where x is independent variable(input) and y is dependent on x(output) we could ignore address column as it required for our model

```
In [12]: x=df1[['Birth Rate', 'Calling Code',
               'Fertility Rate', 'Infant mortality',
               'Life expectancy', 'Maternal mortality ratio','Latitude','Physicians per
               'Longitude'
               ]]
y=df1['Physicians per thousand']
```

## To Split my dataset into training and test data

```
In [13]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [14]: from sklearn.linear_model import LinearRegression
lr= LinearRegression()
lr.fit(x_train,y_train)
```

Out[14]: LinearRegression()

```
In [15]: lr.intercept_
```

Out[15]: -2.531308496145357e-14


```
In [16]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[16]:

	Co-efficient
Birth Rate	-1.762563e-15
Calling Code	1.752775e-17
Fertility Rate	1.289564e-14
Infant mortality	1.704309e-16
Life expectancy	1.996199e-16
Maternal mortality ratio	9.188721e-18
Latitude	-4.783119e-17
Physicians per thousand	1.000000e+00
Longitude	-4.271240e-18

```
In [17]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[17]: <matplotlib.collections.PathCollection at 0x1b5a44a7160>





## Accuracy

```
In [18]: lr.score(x_test,y_test)
```

```
Out[18]: 1.0
```

```
In [19]: lr.score(x_train,y_train)
```

```
Out[19]: 1.0
```

```
In [20]: from sklearn.linear_model import Ridge,Lasso
```

```
In [21]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[21]: Ridge(alpha=10)
```

```
In [22]: rr.score(x_test,y_test)
```

```
Out[22]: 0.9947727105233788
```

```
In [23]: la =Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[23]: Lasso(alpha=10)
```

```
In [24]: la.score(x_test,y_test)
```

```
Out[24]: 0.3593921547432203
```

## ElasticNet

```
In [25]: from sklearn.linear_model import ElasticNet
en = ElasticNet()
en.fit(x_train,y_train)
```

```
Out[25]: ElasticNet()
```

```
In [26]: print(en.coef_)
```

```
[-3.27421153e-02  6.00299313e-05 -0.00000000e+00 -1.68899798e-02
 0.00000000e+00  3.06774698e-06  1.52597188e-03  3.87425791e-01
-7.30657009e-04]
```

```
In [27]: print(en.intercept_)
```

```
2.1209864282200077
```

```
In [28]: print(en.predict(x_test))
```

```
[ 0.90933623  1.25937394  3.25945677  1.72304163  2.91248742  4.28695572
 2.94917222  3.35942739  3.47754088  2.33725144  2.87836436  2.89229255
 2.87829374  1.23539902  0.39479642  2.78442559  0.61122873  1.74481716
-0.1468399   2.76140535  2.19833086  0.71336975  2.6189586   0.42998079
 3.21240503  1.67219904  1.35043644  2.76117253  2.61395252  3.96355486
 0.26986623  2.39005891  1.55907402]
```

```
In [29]: print(en.score(x_test,y_test))
```

```
0.831506599003479
```

## Evaluation Metrics

```
In [30]: from sklearn import metrics
```

```
In [31]: print("Mean Absolute Error",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolute Error 5.374436165466448e-15
```

```
In [32]: print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

```
Mean Squared Error: 4.0151772251877673e-29
```

```
In [33]: print("Root Mean Absolute Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
Root Mean Absolute Error: 6.336542610278706e-15
```